

Who Controls AGI? OpenAI, Microsoft & Governance Tensions

As companies inch towards more advanced AI, a critical question arises: **who will control an AGI, and how will it be governed?** A telling case study is the evolving relationship – and recent friction – between OpenAI and Microsoft, which reveals competing interests in AGI development and foreshadows governance challenges.

OpenAI, creator of ChatGPT and GPT-4, has a unique structure: a capped-profit company governed by a non-profit board, with a mission to develop AGI and ensure it benefits humanity. Microsoft invested heavily in OpenAI (over \$13 billion) and became its primary cloud provider and partner, integrating GPT technology into products like Bing. However, their partnership included a now-famous contractual provision often dubbed “**The AGI Clause.**” According to reports, the OpenAI-Microsoft deal stipulated that if OpenAI succeeds in creating **AGI**, several things happen:

1. **Microsoft loses access to OpenAI’s newest models** – The contract *originally* ran until 2030, but *if* OpenAI’s board *declares* one of its AI models to have achieved AGI, then Microsoft’s exclusive licensing of new OpenAI tech is void. In essence, OpenAI could cut off Microsoft from using or knowing about the AGI; Microsoft would be stuck with older models.
2. **Profit-sharing adjusts** – OpenAI’s board also must determine if this AGI is “sufficient AGI,” meaning it can generate enough future profits (>\$100 billion) to satisfy Microsoft and other investors’ returns. If so, OpenAI can exit Microsoft’s investment with minimal penalty, having proven the value.
3. **Microsoft can’t develop its own AGI** – Critically, as long as the contract is in effect, Microsoft is reportedly *barred from developing AGI on its own*. Microsoft agreed not to pursue a rival AGI independently (perhaps focusing only on applied AI using OpenAI’s tech).

This “AGI Clause” was extraordinary – it basically gave OpenAI’s board the power to **declare an AI model to be “AGI” (defined in OpenAI’s charter as “a highly autonomous system that outperforms humans at most economically valuable work”)** and thereby seize *control* of the technology’s release. Microsoft’s incentive in agreeing to it was presumably the belief that AGI was far-off or unlikely before 2030. Indeed, Microsoft’s CEO Satya Nadella in 2023 downplayed the immediacy of superintelligence, implying “*if that happens, all bets are off*” and it’s a problem to worry about later. Microsoft’s team “*didn’t believe AGI was anywhere close... certainly not before 2030 when the contract expires*”, regarding OpenAI’s insistence on the clause with a bit of an eye-roll.

Fast forward to 2023-2024: **the landscape changed**. OpenAI's GPT-4 launched to global fanfare, and CEO Sam Altman's statements grew bolder about being on the cusp of AGI. By 2025, Sam Altman openly suggested AGI could be achieved *"this year"*. Microsoft, observing this and the furious pace of AI advancement (and competitors like Google DeepMind chasing similar goals), likely grew concerned that **OpenAI might trigger the AGI Clause early** – perhaps declaring one of its upcoming models as "AGI" under their own definition, which is *fuzzy* and at OpenAI's sole discretion. That would leave Microsoft suddenly cut out, despite having paid billions and integrated AI into its strategy. It would also mean Microsoft, due to the contract's non-compete, couldn't have a replacement ready (they weren't building their own GPT-4 competitor because they weren't allowed to "develop AGI" while partnered). In effect, Microsoft risked being **completely dependent on OpenAI for frontier AI and then getting locked out at the critical moment**. No wonder, as Wired reported, *"Microsoft is worried that OpenAI will make that determination [of AGI] prematurely"*.

Tensions came to a head in late 2023 when OpenAI's board (concerned about moves toward commercialization and perhaps other issues) **fired Sam Altman** – a shocking move that was reversed after an employee outcry and high-level intervention. This governance crisis revealed that OpenAI's board and Microsoft had differing priorities. Microsoft immediately tried to bring Altman and others into a new MS-led AI lab when he was fired, showing they weren't entirely aligned with OpenAI's non-profit board's more cautious stance. Ultimately Altman was reinstated and the board reconstituted, but the saga gave Microsoft leverage to **renegotiate the partnership terms**.

According to reports in 2024, Microsoft pushed for **elimination of the AGI Clause** in the new terms. Essentially, Microsoft wanted to remove OpenAI's unilateral power to declare AGI and cut them off. There were even leaks that OpenAI had quietly **removed the term "AGI" from its corporate charter** and from the Microsoft contract, possibly to reduce ambiguity (or to reassure Microsoft that it wouldn't spring an AGI surprise). In other words, the **definition of AGI itself became a bargaining chip in a corporate negotiation**. This is remarkable: a multi-billion dollar partnership **hinged on what counts as AGI**, highlighting how economically and strategically crucial "AGI" has become, even as a concept.

Beyond OpenAI/Microsoft, other companies are aware of the control issue. Google merged its Brain team and DeepMind into one unit (Google DeepMind) in 2023, partly to not fall behind OpenAI. Whoever attains AGI first – the thinking goes – could have an immense advantage (economically and even militarily), so controlling that technology is paramount. This raises governance questions: Should any *one* company have unilateral control over an AGI? OpenAI's original mission was to prevent that scenario (hence the OpenAI charter's

talk of avoiding “*undue concentration of power*” and sharing benefits of AGI). Yet ironically, OpenAI’s close Microsoft tie-up created a situation of corporate jockeying.

We also see **tensions between profit motives and safety/ethical governance**. OpenAI’s board originally included scientists and was willing to halt a CEO for moving too fast; Microsoft, as an investor, prioritized keeping the AI engine humming commercially. The resolution (Altman back, board changed) arguably signaled a shift toward more **corporate control and faster productization**. Indeed, by valuing OpenAI at \$80-\$90 billion and converting it into a “*capped-profit with public benefit*” model, OpenAI is more of a business now. This led some to worry: will safety principles be compromised in the race? Or will tech firms self-regulate effectively?

The OpenAI-Microsoft story illustrates a broader point: **AGI governance might require new structures**. Ideas floated include multi-stakeholder oversight (perhaps an international AGI project with government, academia, and industry involvement to avoid one-company dominance) and mechanisms like *windfall clauses* (where companies commit to distributing a portion of AGI-generated profits to society). OpenAI actually has a *windfall clause* pledge publicly, and its Charter says if an AGI is developed outside OpenAI they will stop and cooperate. Whether these idealistic commitments hold under competitive pressure is uncertain.

Another dimension is **government vs corporate control**. If a corporation is close to AGI, should governments step in? We’ve seen the U.S. and China each ensure deep involvement – the U.S. via funding and defense interest (the Pentagon is now investing in frontier AI, e.g., the DARPA “AI Forward” initiative), and China via direct state oversight. It’s plausible that an AGI would not remain solely in private hands; it might become akin to nuclear technology – too powerful to be left without government regulation or even international treaties.

In essence, the **question “Who controls AGI?” is as much about corporate power as technical capability**. 2024’s events showed that even defining AGI can trigger contractual escape hatches worth hundreds of billions. As Steven Levy wrote, “*pretty much everything happening in AI can be explained by The Clause*.” It epitomized the split worldview: OpenAI’s leaders **truly believe AGI is near and world-changing**, thus wanting freedom to act, while Microsoft’s leaders were initially skeptical of the timeline and insisted on protection (which now they reconsider seeing rapid progress).

Moving forward, if multiple entities approach AGI, we might see **agreements or consortiums** to jointly govern it (to avoid a dangerous race). Already, Google DeepMind and Anthropic, etc., hold talks on safety standards. If one player appears to lead, others may band together to catch up or to lobby for oversight.

The bottom line is, **the control of AGI will likely not be absolute** – various stakeholders (corporations, governments, international bodies, and the public) will exert influence. The OpenAI/Microsoft saga is a microcosm of that, compressing issues of trust, profit, and principle. It underscores that AGI development is not happening in a social vacuum; it's entangled with corporate deals, legal clauses, and power plays. As AGI draws closer, expect more of these governance dramas: they are how society will hash out *who gets to steer the most powerful technology of our time* and under what constraints.